

Yash Raghuvanshi

Machine Learning Engineer

+91 78782 81860 | tanmaysraghuvanshi23@gmail.com | Kota, Rajasthan | github.com/Tanmay23yash

SUMMARY

Artificial Intelligence Engineer with 1+ year at L&T Technology Services building AI agents, CI/CD pipelines and backend integrations. Builds LLM applications end to end with RAG, vector search and multi-agent orchestration, and has trained transformers from scratch and fine-tuned LLMs with LoRA in PyTorch.

EXPERIENCE

L&T Technology Services (LTTS)

[September 2025] – Present

[Associate Engineer]

India

- Engineered an autonomous AI agent that parses, refactors and modifies codebases from natural-language instructions.
- Integrated agentic developer tooling (Claude Code) into the IDE workflow, speeding up development and refactoring.
- Designed Jenkins pipelines for automated testing and deployment.
- Resolved critical API and backend integration issues with Postman while providing cross-team technical support.

Internshala

[June 2023] – [August 2023]

Machine Learning Intern

Remote

- Built a disease prediction model (Logistic Regression) and compared KNN, SVC and Decision Tree classifiers.

PROJECTS

HireMind – AI Mock Interview Platform

tryhiremind.vercel.app

Python, FastAPI, React, TypeScript, PostgreSQL, pgvector, Gemini API, Jenkins

- Built and deployed a full-stack platform that runs adaptive technical interviews from a candidate's resume and a target job description; React 19 front end on Vercel, FastAPI back end on Render.
- Implemented a RAG pipeline: section-aware resume chunking, 768-dim Gemini embeddings in **pgvector**, and cosine top-k retrieval that grounds each question and grade in the resume.
- Designed a multi-agent orchestrator (planner, question generator, evaluator, report writer) with Pydantic-structured outputs, 5-level adaptive difficulty and job-description skill-gap analysis.
- Added JWT and Google OAuth login, correlation-ID logging with per-step LLM token and latency tracing, and Jenkins CI running 90+ pytest and Vitest tests with SonarQube analysis.

Mini-GPT – Hinglish GPT from Scratch vs LoRA-Tuned LLM

github.com/Tanmay23yash/mini-gpt

PyTorch, Transformers, LoRA, Hugging Face, Gradio

- Implemented a 5.3M-parameter GPT from scratch in PyTorch: byte-level BPE tokenizer, multi-head causal self-attention, pre-norm transformer blocks, AdamW with warmup and cosine decay, top-k/top-p sampling.
- Pre-trained on 98M tokens of Hinglish tweets on a 4 GB GPU in 40 minutes (validation perplexity 8,659 to 150); pre-training raised sentiment macro F1 by 3.9 points (0.640 to 0.679).
- Benchmarked against TF-IDF, BiLSTM/BiGRU and **Qwen3.5-0.8B** (zero-shot, few-shot, **LoRA**) on SemEval-2020 SentiMix; LoRA scored best at 0.714 macro F1 over 3 seeds with bootstrap 95% confidence intervals.
- Deployed a Gradio demo with attention maps: huggingface.co/spaces/Yash230901/mini-gpt-hinglish

TECHNICAL SKILLS

Languages: Python, SQL, C, C++

ML / AI: PyTorch, TensorFlow, scikit-learn, Hugging Face Transformers, LoRA / PEFT, LLMs, RAG, Embeddings, Vector Search, Prompt Engineering, AI Agents, NLP, Deep Learning, Gemini API, Pandas, NumPy

Backend & Data: FastAPI, REST APIs, Pydantic, Django, PostgreSQL, pgvector, SQLAlchemy

DevOps & Tools: Git, GitHub, Jenkins, CI/CD, SonarQube, Postman, pytest, Render, Vercel, Claude Code

EDUCATION

Amity University, Gwalior

Aug 2025

B.Tech in Computer Science and Engineering – CGPA: 8.44/10

Gwalior, Madhya Pradesh

CERTIFICATIONS & ACHIEVEMENTS

- Oracle Cloud Infrastructure 2025 Certified Generative AI Professional; OCI 2025 Certified Foundations Associate
- Machine Learning and Deep Learning (Coursera / Stanford University); EF SET English Certificate (C2 Proficient)
- CodeChef: highest rating 1307, Gold badge (contests), Bronze badge (problem solving) – handle singh_tanmay23